

432 Class 02

<https://thomaseLove.github.io/432-2023/>

2023-01-19

Today's Agenda

- ① Comparing Means
- ② Comparing Rates
- ③ Fitting Linear Models
- ④ Setting Up Lab 1, due Monday 2023-01-23 at 9 PM.

The most relevant sections of the Course Notes are Chapters 1-5.

Today's R Setup

```
knitr::opts_chunk$set(comment = NA)

library(broom) # for tidy, glance and augment
library(car) # for vif
library(Epi) # for twoby2
library(GGally) # for ggpairs
library(knitr) # for kable
library(MKinfer) # for boot.t.test
library(mosaic) # for favstats
library(naniar) # deal with missing values
library(nhanesA) # source of data
library(vcd) # for mosaic (plot) and assoc (plot)
library(janitor) # for tabyl and other things
library(tidyverse) # for all kinds of things

theme_set(theme_bw())
```

Section 1

Building an NHANES Data Set (see Course Notes
Chapters 1-2)

How I Built Our Data (2017 - March 2020 NHANES)

Variables from P-DEMO:

- SEQN
- RIDAGEYR (age) *restricted to ages 26-42 here*
- DMDEDUC2 (educ)

Variables from BPXO (linked by SEQN):

- BPXOSY1 (sbp1)
- BPXOSY2 (sbp2)
- BPXOSY3 (sbp3)

Variables from HUQ (linked by SEQN)

- HUQ010 (sroh)
- HUQ071 (hospital)
- HUQ090 (mentalh)

Total: 1982 observations on 9 variables: includes all available NHANES participants ages 26-42 with complete data on these nine variables.

Building the Data (using nhanesA package)

```
p_demo <- nhanes('P_DEMO') |>
  select(SEQN, RIDAGEYR, DMDEDUC2)

p_bpxo <- nhanes('P_BPXO') |>
  select(SEQN, BPXOSY1, BPXOSY2, BPXOSY3)

p_huq <- nhanes('P_HUQ') |>
  select(SEQN, HUQ010, HUQ071, HUQ090)

df_list <- list(p_demo, p_bpxo, p_huq)

nh_raw <- df_list |>
  reduce(left_join, by = 'SEQN') |>
  drop_na() |>
  filter(RIDAGEYR >= 26 & RIDAGEYR <= 42) |>
  as_tibble()
```

Renaming and Cleaning Variables (1)

```
nh1982 <- nh_raw |>
  rename(age = RIDAGEYR, educ = DMDEDUC2,
         sbp1 = BPXOSY1, sbp2 = BPXOSY2,
         sbp3 = BPXOSY3, sroh = HUQ010,
         hospital = HUQ071, mentalh = HUQ090) |>
  replace_with_na_at(
    .vars = c("educ", "sroh", "hospital", "mentalh"),
    condition = ~ .x %in% c(7,9)) |>
  mutate(across(c(hospital, mentalh), ~ 2 - .x)) |>
  mutate(mean_sbp = (sbp1 + sbp2 + sbp3)/3,
         SEQN = as.character(SEQN))
```

Renaming and Cleaning Variables (2)

```
nh1982 <- nh1982 |>
  mutate(educ = fct_recode(factor(educ),
    "Less than 9th Grade" = "1",
    "9th - 11th Grade" = "2",
    "High School Grad" = "3",
    "Some College / AA" = "4",
    "College Grad" = "5")) |>
  mutate(sroh = fct_recode(factor(sroh),
    "Excellent" = "1",
    "Very Good" = "2",
    "Good" = "3",
    "Fair" = "4",
    "Poor" = "5")) |>
  drop_na()

write_rds(nh1982, "c02/data/nh1982.Rds")
```


nh1982

```
glimpse(nh1982)
```

Rows: 1,982

Columns: 10

```
$ SEQN      <chr> "109266", "109273", "109291", "109297", "1093
$ age       <dbl> 29, 36, 42, 30, 30, 28, 33, 41, 35, 30, 41, 3
$ educ      <fct> College Grad, Some College / AA, College Grad
$ sbp1      <dbl> 99, 116, 107, 105, 118, 110, 110, 106, 162, 1
$ sbp2      <dbl> 99, 110, 111, 105, 123, 110, 105, 107, 148, 1
$ sbp3      <dbl> 99, 115, 107, 102, 125, 110, 108, 113, 163, 1
$ sroh      <fct> Good, Good, Fair, Very Good, Good, Very Good,
$ hospital  <dbl> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
$ mentalh   <dbl> 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0,
$ mean_sbp  <dbl> 99.00000, 113.66667, 108.33333, 104.00000, 12
```

Codebook (excerpt, without SEQN)

Variable	Description (n = 1982)
age	Age in years (range 26-42, mean = 34)
meansbp	Mean of sbp1, sbp2, sbp3 in mm Hg (range: 76 to 209, mean 116): we'll also use sbp1, sbp2 and sbp3.
hospital	1 if hospitalized in last 12m, else 0 (8% are 1)
mentalh	1 if saw a mental health professional in last 12m, else 0 (12% are 1)
sroh	Self-reported Overall Health (5 levels: see next slide)
educ	Educational Attainment (5 levels: see next slide)

SROH and Educational Attainment

```
nh1982 |> tabyl(sroh) |> adorn_pct_formatting()
```

	sroh	n	percent
Excellent	294	14.8%	
Very Good	598	30.2%	
Good	728	36.7%	
Fair	321	16.2%	
Poor	41	2.1%	

```
nh1982 |> tabyl(educ) |> adorn_pct_formatting()
```

	educ	n	percent
Less than 9th Grade	90	4.5%	
9th - 11th Grade	209	10.5%	
High School Grad	418	21.1%	
Some College / AA	677	34.2%	
College Grad	588	29.7%	

Ingesting the Data (from .Rds)

If you don't want to work through the `nhanesA` import and tidying, you can simply work with the `nh1982.Rds` file provided on our 432-data page.

```
nh1982 <- read_rds("c02/data/nh1982.Rds")  
  
## not run here...
```

Section 2

Comparing Means (see Course Notes Chapter 3)

Paired or Independent Samples?

In Analysis 1, we will compare the means of SBP1 and SBP2 for our 1982 participants.

In Analysis 2, we will compare the mean of SBP2 between our 159 participants who were hospitalized and the 1823 who were not?

Which of these analyses uses paired samples, and why?

Paired Samples Analysis


```
nh1982 <- nh1982 |> mutate(SBP_diff = sbp1 - sbp2)

favstats(~ SBP_diff, data = nh1982)
```

min	Q1	median	Q3	max	mean	sd	n	missing
-27	-3	0	3	37	0.2482341	5.279749	1982	0

```
ggplot(nh1982, aes(sample = SBP_diff)) +
  geom_qq() + geom_qq_line(col = "red") +
  labs(title = "Normal Q-Q plot",
       subtitle = "SBP1 - SBP2 Differences",
       x = "Standard Normal Distribution",
       y = "Observed SBP1-SBP2 values")
```

Normal Q-Q plot of Paired SBP Differences



Comparing Paired Samples

Want a 90% confidence interval for the true mean of the paired SBP1 - SBP2 differences:

- t-based approach (equivalent to linear model) assumes Normality
- Wilcoxon signed rank approach doesn't assume Normality but makes inferences about the pseudo-median, not the mean
- bootstrap doesn't assume Normality, and describes the mean

```
set.seed(20230117)
boot.t.test(nh1982$SBP_diff, conf.level = 0.9,
            boot = TRUE, R = 999)
```

Results on the next slide...

Bootstrap for Mean of SBP1-SBP2 Differences

Bootstrap One Sample t-test

```
data:  nh1982$SBP_diff
bootstrap p-value = 0.05205
bootstrap mean of x (SE) = 0.2555684 (0.1184953)
90 percent bootstrap percentile confidence interval:
  0.06609485 0.44611504
```

Results without bootstrap:

```
t = 2.0931, df = 1981, p-value = 0.03646
alternative hypothesis: true mean is not equal to 0
90 percent confidence interval:
  0.05307362 0.44339459
sample estimates:
mean of x
0.2482341
```

Comparing sbp3 by hospital: Independent Samples


```
favstats(sbp3 ~ hospital, data = nh1982) |>  
  select(-missing) |>  
  kable(digits = 2)
```

hospital	min	Q1	median	Q3	max	mean	sd	n
0	60	107	115	124.0	204	116.11	14.51	1823
1	88	104	113	124.5	195	116.71	18.50	159

```
ggplot(nh1982, aes(x = factor(hospital), y = sbp3)) +  
  geom_violin(aes(fill = factor(hospital))) +  
  geom_boxplot(width = 0.3, notch = TRUE) +  
  guides(fill = "none") +  
  labs(title = "SBP (3rd reading) by Hospitalization")
```

SBP (3rd reading) vs. Hospitalization Status

SBP (3rd reading) by Hospitalization



Two Independent Samples, Comparing Means

Want a 90% confidence interval for the difference in means of SBP3 for people who were hospitalized - those who were not.

- Pooled t-based approach (equivalent to linear model) assumes Normality and equal population variances
- Welch t-based approach assumes Normality only
- bootstrap assumes neither
- Wilcoxon-Mann-Whitney rank sum assumes neither, but assesses a difference in locations, not the mean

Pooled t test approach via linear model

```
lm2 <- lm(sbp3 ~ hospital, data = nh1982)

tidy(lm2, conf.int = TRUE, conf.level = 0.90) |>
  kable(digits = 2)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	116.11	0.35	333.54	0.00	115.54	116.69
hospital	0.60	1.23	0.49	0.63	-1.42	2.62

```
glance(lm2) |> select(r.squared, sigma) |>
  kable(digits = c(5,2))
```

r.squared	sigma
0.00012	14.86

Section 3

Comparing Rates (see Course Notes, Chapter 4)

A Two-by-Two Contingency Table

```
nh1982 |> tabyl(mentalh, hospital) |>
  adorn_totals(where = c("row", "col")) |>
  adorn_title()
```

	hospital		
mentalh	0	1	Total
0	1613	122	1735
1	210	37	247
Total	1823	159	1982

Standard Epidemiological Format

```
nh1982 <- nh1982 |>
  mutate(mental_h_f = fct_recode(factor(mental_h),
    "Saw MHP" = "1", "No MHP" = "0"),
    mental_h_f = fct_relevel(mental_h_f,
      "Saw MHP", "No MHP"),
    hospital_f = fct_recode(factor(hospital),
      "Hosp." = "1", "No Hosp." = "0"),
    hospital_f = fct_relevel(hospital_f,
      "Hosp.", "No Hosp.))

nh1982 |> tabyl(mental_h_f, hospital_f)
```

mental_h_f	Hosp.	No Hosp.
Saw MHP	37	210
No MHP	122	1613

Two by Two Table Analysis

```
twoby2(nh1982$mentalh_f, nh1982$hospital_f, conf.level = 0.90)
```

2 by 2 table analysis:

Outcome : Hosp.

Comparing : Saw MHP vs. No MHP

	Hosp.	No Hosp.	P(Hosp.)	90% conf. interval	
Saw MHP	37	210	0.1498	0.1161	0.1911
No MHP	122	1613	0.0703	0.0609	0.0811

	90% conf. interval		
Relative Risk:	2.1303	1.5977	2.8405
Sample Odds Ratio:	2.3295	1.6723	3.2449
Conditional MLE Odds Ratio:	2.3282	1.6287	3.2894
Probability difference:	0.0795	0.0442	0.1217

A Larger Two-Way Table

What is the association of Educational Attainment with Self-Reported Overall Health?

```
nh1982 |> tabyl(educ, sroh) |>  
  adorn_totals(where =c("row","col"))|> adorn_title()
```

	sroh						
educ	Excellent	Very Good	Good	Fair	Poor	Total	
Less than 9th Grade	10	7	36	33	4	90	
9th - 11th Grade	21	40	81	59	8	209	
High School Grad	50	94	168	98	8	418	
Some College / AA	72	220	264	104	17	677	
College Grad	141	237	179	27	4	588	
Total	294	598	728	321	41	1982	


Our 5x5 Table, showing SROH Proportions

```
nh1982 |> tabyl(educ, sroh) |>  
  adorn_totals(where = c("row")) |>  
  adorn_percentages(denominator = "row") |>  
  adorn_pct_formatting() |> adorn_title()
```

	sroh				
educ	Excellent	Very Good	Good	Fair	Poor
Less than 9th Grade	11.1%	7.8%	40.0%	36.7%	4.4%
9th - 11th Grade	10.0%	19.1%	38.8%	28.2%	3.8%
High School Grad	12.0%	22.5%	40.2%	23.4%	1.9%
Some College / AA	10.6%	32.5%	39.0%	15.4%	2.5%
College Grad	24.0%	40.3%	30.4%	4.6%	0.7%
Total	14.8%	30.2%	36.7%	16.2%	2.1%

Mosaic Plot for our 5x5 Table

```
mosaic(~ educ + sroh, data = nh1982, highlighting = "sroh")
```



Pearson χ^2 test for our 5x5 Table


```
chisq.test(xtabs(~ educ + sroh, data = nh1982))
```

Pearson's Chi-squared test

```
data:  xtabs(~educ + sroh, data = nh1982)  
X-squared = 225.99, df = 16, p-value < 2.2e-16
```

Association Plot for our 5x5 Table

```
assoc(~ educ + sroh, data = nh1982)
```



Section 4

Fitting Linear Models (see Course Notes, Chapter 5)

We'll fit two models today


- 1 Predict mean SBP using Age alone.
- 2 Predict mean SBP (across three readings) using Age, Self-Reported Overall Health Status and Hospitalization Status.

```
temp_mod1 <- lm(mean_sbp ~ age, data = nh1982)
temp_mod2 <- lm(mean_sbp ~ age + sroh + hospital,
                 data = nh1982)
```

Note that I'm not doing any predictive validation today (remember that I did that in Class 1), so I won't split the sample.

Box-Cox Plot to suggest potential outcome transformations

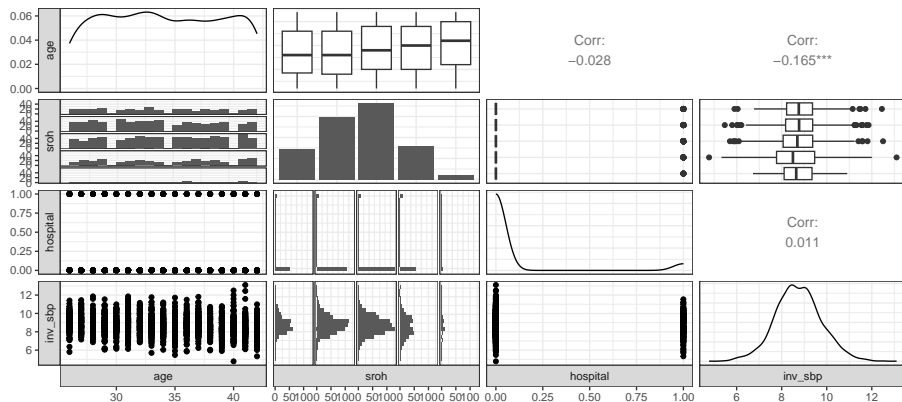
```
boxCox(temp_mod2)
```



```
nh1982 <- nh1982 |> mutate(inv_sbp = 1000/mean_sbp)
```

Scatterplot Matrix (from ggpairs())

```
ggpairs(nh1982, columns = c(2, 7, 8, 14), switch = "both",  
        lower=list(combo=wrap("facethist", bins=20)))
```



Checking Collinearity: Variance Inflation Factors

```
vif(lm(inv_sbp ~ age + sroh + hospital, data = nh1982))
```

	GVIF	Df	$GVIF^{(1/(2*Df))}$
age	1.008723	1	1.004352
sroh	1.020544	4	1.002545
hospital	1.013797	1	1.006875

Tidied Coefficients for Model m1

```
m1 <- lm(inv_sbp ~ age, data = nh1982)

tidy(m1, conf.int = TRUE, conf.level = 0.9) |>
  kable(digits = 2)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	9.93	0.16	61.52	0	9.66	10.20
age	-0.03	0.00	-7.44	0	-0.04	-0.03

Tidied Coefficients for Model m2

```
m2 <- lm(inv_sbp ~ age + sroh + hospital, data = nh1982)

tidy(m2, conf.int = TRUE, conf.level = 0.9) |>
  kable(digits = 2)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	9.99	0.17	58.76	0.00	9.71	10.27
age	-0.03	0.00	-7.19	0.00	-0.04	-0.03
srohVery Good	-0.06	0.07	-0.76	0.45	-0.17	0.06
srohGood	-0.11	0.07	-1.56	0.12	-0.23	0.01
srohFair	-0.27	0.08	-3.21	0.00	-0.40	-0.13
srohPoor	-0.18	0.17	-1.03	0.30	-0.46	0.10
hospital	0.05	0.08	0.55	0.58	-0.09	0.19





Fit Summaries for Models m1 and m2

```
bind_rows(glance(m1), glance(m2)) |>
  mutate(model = c("m1", "m2")) |>
  select(model, r2 = r.squared, adjr2 = adj.r.squared,
         sigma, AIC, BIC, nobs, df, df.residual) |>
  kable(digits = c(0, 3, 3, 3, 1, 1, 0, 0))
```

model	r2	adjr2	sigma	AIC	BIC	nobs	df	df.residual
m1	0.027	0.027	1.022	5713.8	5730.6	1982	1	1980
m2	0.033	0.030	1.020	5711.2	5755.9	1982	6	1975

Which model appears to fit the data better?

Residual Plots for Model m2



Making a Prediction in New Data

Suppose a new person is age 29, was not hospitalized, and their SROH is "Good". What is their predicted mean systolic blood pressure?

- Our models predict $1000/\text{mean_sbp}$ and augment places that prediction into `.fitted`.
- To invert, divide `.fitted` by 1000, then take the reciprocal of that result. That's just $1000/.\text{fitted}$.

```
new_person <- tibble(age = 29, sroh = "Good", hospital = 0)
bind_rows(augment(m1, newdata = new_person),
          augment(m2, newdata = new_person)) |>
  mutate(model = c("m1", "m2"), fit_meansbp = 1000/.\fitted) |>
  select(model, fit_meansbp, .fitted, age, sroh, hospital) |>
```

model	fit_meansbp	.fitted	age	sroh	hospital
m1	112.114	8.920	29	Good	0
m2	112.309	8.904	29	Good	0

Section 5

Setting Up Lab 1, due Monday 2023-01-23 at 9 PM

Lab 1 Question 1

I provide some County Health Rankings data for Ohio's 88 counties. You create a visualization involving information from at least three different variables using R and Quarto.

- Include proper labels and a meaningful title.
- Include a caption (75 words or fewer) that highlights the key result.
- What is the question you are trying to answer with this visualization?

There is a Quarto template for Lab 1, in addition to the data set.

Lab 1 Question 2

Create a linear regression model to predict `obese_pct` as a function of `food_env` and `median_income` (all of these are quantitative variables.)

- a. Specify and fit the model, interpret `food_env` coefficient.
- b. Evaluate quality of model in terms of adherence to regression assumptions via four key residual plots.
- c. Build a nice table comparing your model to a simple regression for `obese_pct` using only `food_env`, and then reflect on your findings.

Next Week?

- Lab 1 due Monday 9 PM (Answer Sketch available Tuesday)
- Developing Inferences Using Survey Weights
- Linear Regression and ANOVA/ANCOVA models