

432 Class 17

<https://thomaseLove.github.io/432-2023/>

2023-03-07

Today's Topic

Regression Models for Ordered Multi-Categorical Outcomes

- Applying to Graduate School: A First Example
- Proportional Odds Logistic Regression Models
 - Using `polr`
 - Using `lrm`
- Understanding and Interpreting the Model
- Testing the Proportional Odds Assumption
- Picturing the Model Fit

Chapter 27 of the Course Notes describes this material.

Setup

```
knitr::opts_chunk$set(comment=NA)
options(width = 60)

library(janitor)
library(GGally)
library(scales)
library(knitr)
library(rms)
library(MASS)
library(nnet)
library(tidyverse)

theme_set(theme_bw())
```

Section 1

Applying to Graduate School

These are **simulated** data

This is a simulated data set of 530 students, looking at factors that influence the decision of whether to apply to graduate school.

College juniors are asked if they are unlikely, somewhat likely, or very likely to apply to graduate school. Hence, our outcome variable has three categories. Data on parental educational status, whether the undergraduate institution is public or private, and current GPA is also collected. The researchers have reason to believe that the “distances” between these three points are not equal. For example, the “distance” between “unlikely” and “somewhat likely” may be shorter than the distance between “somewhat likely” and “very likely”.

```
gradschool <-  
  read_csv("c17/data/gradschool_new.csv",  
           show_col_types = FALSE) |>  
  type.convert(as.is = FALSE) |>  
  clean_names()
```

The gradschool data and my **Source**

The **gradschool** example is adapted from [this UCLA site](#).

- There, they look at 400 students.
- I simulated a new data set containing 530 students.

Variable	Description
student	subject identifying code (A001 - A530)
apply	3-level ordered outcome: “unlikely”, “somewhat likely” and “very likely” to apply
pared	1 = at least one parent has a graduate degree, else 0
public	1 = undergraduate institution is public, else 0
gpa	student’s undergraduate grade point average (max 4.00)

Ensuring that our outcome is an ordered factor

```
gradschool <- gradschool |>
  mutate(apply = fct_relevel(apply, "unlikely",
                             "somewhat likely", "very likely"),
         apply = factor(apply, ordered = TRUE))

is.ordered(gradschool$apply)
```

```
[1] TRUE
```

The gradschool tibble

```
gradschool
```

```
# A tibble: 530 x 5
```

	student	apply	pared	public	gpa
	<fct>	<ord>	<int>	<int>	<dbl>
1	A001	very likely	0	0	3.41
2	A002	unlikely	0	0	2.38
3	A003	somewhat likely	0	0	3.35
4	A004	unlikely	0	1	3.45
5	A005	unlikely	1	1	3.27
6	A006	somewhat likely	1	0	3.41
7	A007	somewhat likely	0	0	2.83
8	A008	unlikely	0	0	3.64
9	A009	unlikely	0	0	2.52
10	A010	unlikely	0	0	2.36

```
# ... with 520 more rows
```


Numerical Description of gradschool data

```
> describe(gradschool)
gradschool
```

5 Variables			530 Observations				

student							
n	missing	distinct					
530	0	530					
lowest : A001 A002 A003 A004 A005, highest: A526 A527 A528 A529 A530							

apply							
n	missing	distinct					
530	0	3					
Value							
somewhat likely	unlikely	very likely					
Frequency	172	303					
Proportion	0.325	0.572	0.104				

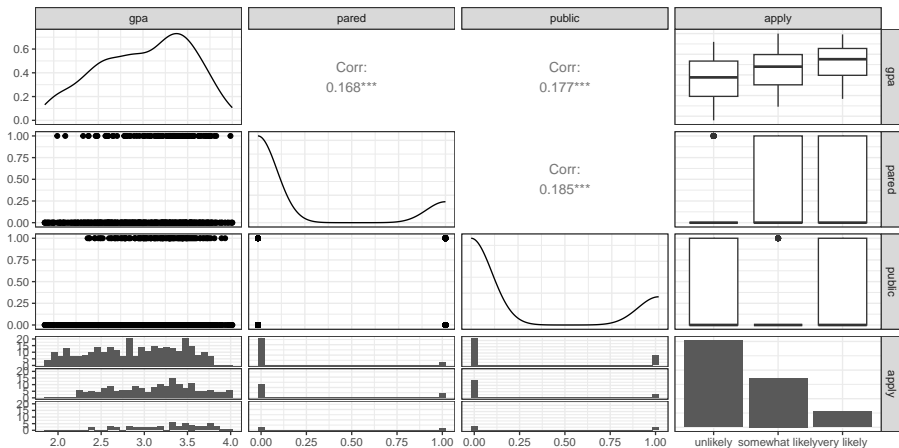
pared							
n	missing	distinct	Info	Sum	Mean	Gmd	
530	0	2	0.47	103	0.1943	0.3137	

public							
n	missing	distinct	Info	Sum	Mean	Gmd	
530	0	2	0.555	130	0.2453	0.3709	

gpa							
n	missing	distinct	Info	Mean	Gmd	.05	.10
530	0	186	1	3.015	0.5919	2.104	2.279
.25	.50	.75	.90	.95			
2.610	3.080	3.440	3.660	3.760			
lowest : 1.90 1.91 1.92 1.93 1.94, highest: 3.95 3.97 3.98 3.99 4.00							

Scatterplot Matrix (run with #| message: FALSE)


```
ggpairs(gradschool |> select(gpa, pared, public, apply))
```



Bar Chart of apply classifications with %s (code)

```
ggplot(gradschool, aes(x = apply, fill = apply)) +  
  geom_bar(aes(y =  
    (after_stat(count)/sum(after_stat(count))))) +  
  geom_text(aes(y =  
    (after_stat(count))/sum(after_stat(count)),  
    label = scales::percent((after_stat(count)) /  
      sum(after_stat(count))),  
    stat = "count", vjust = 1.5,  
    color = "white", size = 5) +  
  scale_y_continuous(labels = scales::percent) +  
  scale_fill_brewer(palette = "Set1") +  
  guides(fill = "none") +  
  labs(y = "Percentage")
```

Bar Chart of apply classifications with %s (result)



Data (besides gpa) as Cross-Tabulation


```
fctable(xtabs(~ public + apply + pared, data = gradschool))
```

		pared	0	1
public	apply			
0	unlikely		206	17
	somewhat likely		111	32
	very likely		22	12
1	unlikely		62	18
	somewhat likely		15	14
	very likely		11	10

apply percentages by public, pared (code)

```
ggplot(gradschool, aes(x = apply, fill = apply)) +  
  geom_bar() +  
  scale_fill_brewer(palette = "Set1") +  
  guides(fill = "none") +  
  facet_grid(pared ~ public, labeller = "label_both")
```


apply percentages by public, pared



Breakdown of gpa by apply (code)

```
ggplot(gradschool, aes(x = apply, y = gpa, fill = apply)) +  
  geom_violin(trim = TRUE) +  
  geom_boxplot(col = "white", width = 0.2) +  
  scale_fill_brewer(palette = "Set1") +  
  guides(fill = "none")
```



Breakdown of gpa by apply



gpa by three other variables (code)

```
ggplot(gradschool, aes(x = apply, y = gpa)) +  
  geom_boxplot(aes(fill = apply), size = .75) +  
  geom_jitter(alpha = .25) +  
  facet_grid(pared ~ public, margins = TRUE,  
             labeller = "label_both") +  
  scale_fill_brewer(palette = "Set1") +  
  guides(fill = "none") +  
  theme(axis.text.x =  
        element_text(angle = 45, hjust = 1, vjust = 1))
```

gpa by three other variables



Section 2

Proportional Odds Logit Model via `polr`

Fitting the POLR model with MASS::polr

We use the `polr` function from the MASS package:

```
mod_p1 <- polr(apply ~ pared + public + gpa,  
              data = gradschool, Hess=TRUE)
```

The `polr` name comes from proportional odds logistic regression, highlighting a key assumption of this model.

`polr` uses the standard formula interface in R for specifying a regression model with outcome followed by predictors. We also specify `Hess=TRUE` to have the model return the observed information matrix from optimization (called the Hessian) which is used to get standard errors.

Obtaining Predicted Probabilities from `mod_p1`

To start we'll obtain predicted probabilities, which are usually the best way to understand the model.

For example, we can vary `gpa` for each level of `pared` and `public` and calculate the model's estimated probability of being in each category of `apply`.

First, create a new tibble of values to use for prediction.

```
newdat <- tibble(  
  pared = rep(0:1, 200),  
  public = rep(0:1, each = 200),  
  gpa = rep(seq(from = 1.9, to = 4, length.out = 100), 4))
```

Obtaining Predicted Probabilities from mod_p1

Now, make predictions using model mod_p1:

```
newdat_p1 <- cbind(newdat,  
                    predict(mod_p1, newdat, type = "probs"))  
head(newdat_p1) |> kable(digits = 3)
```

pared	public	gpa	unlikely	somewhat likely	very likely
0	0	1.900	0.846	0.132	0.022
1	0	1.921	0.629	0.302	0.069
0	0	1.942	0.840	0.137	0.024
1	0	1.964	0.617	0.310	0.073
0	0	1.985	0.833	0.142	0.025
1	0	2.006	0.606	0.318	0.076

Reshape data

Now, we reshape the data with `pivot_longer`:

```
newdat_long <-  
  pivot_longer(newdat_p1,  
               cols = c("unlikely":"very likely"),  
               names_to = "level",  
               values_to = "probability") |>  
  mutate(level = fct_relevel(level, "unlikely",  
                             "somewhat likely"))
```

Result on next slide...

The newdat_long data

```
# A tibble: 1,200 x 5
```


	pared	public	gpa	level	probability
	<int>	<int>	<dbl>	<fct>	<dbl>
1	0	0	1.9	unlikely	0.846
2	0	0	1.9	somewhat likely	0.132
3	0	0	1.9	very likely	0.0225
4	1	0	1.92	unlikely	0.629
5	1	0	1.92	somewhat likely	0.302
6	1	0	1.92	very likely	0.0694
7	0	0	1.94	unlikely	0.840
8	0	0	1.94	somewhat likely	0.137
9	0	0	1.94	very likely	0.0236
10	1	0	1.96	unlikely	0.617

```
# ... with 1,190 more rows
```

Plot the prediction results... (code)

```
ggplot(newdat_long, aes(x = gpa, y = probability,  
                        color = level)) +  
  geom_line(size = 1.5) +  
  scale_color_brewer(palette = "Set1") +  
  facet_grid(pared ~ public, labeller="label_both")
```

Plot the prediction results...



Cross-Tabulation of Predicted/Observed Classifications

Predictions in the rows, Observed in the columns

```
addmargins(table(predict(mod_p1), gradschool$apply))
```

	unlikely	somewhat	likely	very likely	Sum
unlikely	264		112	29	405
somewhat likely	39		60	25	124
very likely	0		0	1	1
Sum	303		172	55	530

We only predict one subject to be in the “very likely” group by modal prediction.

Describing the Proportional Odds Logistic Model

Our outcome, `apply`, has three levels. Our model has two logit equations:

- one estimating the log odds that `apply` will be less than or equal to 1 (`apply` = “unlikely”)
- one estimating the log odds that `apply` ≤ 2 (`apply` = “unlikely” or “somewhat likely”)

That’s all we need to estimate the three categories, since $\Pr(\text{apply} \leq 3) = 1$, because “very likely” is the maximum category for `apply`.

Parameters of the POLR Model

- The parameters to be fit include two intercepts:
 - ζ_1 will be the unlikely|somewhat likely parameter
 - ζ_2 will be the somewhat likely|very likely parameter

We'll have a total of five free parameters when we add in the slopes (β) for pared, public and gpa.

- The two logistic equations that will be fit differ only by their intercepts.

summary(mod_p1)

Call:

```
polr(formula = apply ~ pared + public + gpa, data = gradschool,
      Hess = TRUE)
```

Coefficients:

	Value	Std. Error	t value
pared	1.1525	0.2184	5.276
public	-0.4949	0.2195	-2.254
gpa	1.1416	0.1850	6.171

Intercepts:

	Value	Std. Error	t value
unlikely somewhat likely	3.8727	0.5721	6.7692
somewhat likely very likely	5.9413	0.6063	9.7993

Residual Deviance: 900.9629

AIC: 910.9629

Understanding the Model

$$\text{logit}[Pr(\text{apply} \leq 1)] = \zeta_1 - \beta_1 \text{pared} - \beta_2 \text{public} - \beta_3 \text{gpa}$$

$$\text{logit}[Pr(\text{apply} \leq 2)] = \zeta_2 - \beta_1 \text{pared} - \beta_2 \text{public} - \beta_3 \text{gpa}$$

in general. In our setting, we have ...

The mod_p1 equations...

$$\text{logit}[Pr(\text{apply} \leq \text{unlikely})] = 3.87 - 1.15\text{pared} - (-0.49)\text{public} - 1.14\text{gpa}$$

and

$$\text{logit}[Pr(\text{apply} \leq \text{somewhat})] = 5.94 - 1.15\text{pared} - (-0.49)\text{public} - 1.14\text{gpa}$$

confint(mod_p1)

Confidence intervals for the slope coefficients on the log odds scale can be estimated in the usual way.

	2.5 %	97.5 %
pared	0.7257019	1.58305735
public	-0.9320573	-0.07029727
gpa	0.7837559	1.50974002

These CIs describe results in units of ordered log odds.

- For example, for a one unit increase in gpa, we expect a 1.14 increase in the expected value of apply (95% CI 0.78, 1.51) in the log odds scale, holding pared and public constant.
- This would be more straightforward if we exponentiated.

Exponentiating the Coefficients

```
exp(coef(mod_p1))
```

	pared	public	gpa
	3.1660446	0.6096623	3.1318247

```
exp(confint(mod_p1))
```

		2.5 %	97.5 %
pared	2.0661808	4.8698218	
public	0.3937428	0.9321167	
gpa	2.1896811	4.5255541	

Interpreting the Exponentiated Coefficients

Variable	Estimate	95% CI
gpa	3.13	(2.19, 4.53)
public	0.61	(0.39, 0.93)
pared	3.17	(2.07, 4.87)

- When a student's gpa increases by 1 unit, the odds of moving from “unlikely” applying to “somewhat likely” or “very likely” applying are multiplied by 3.13 (95% CI 2.19, 4.52), all else held constant.
- For public, the odds of moving from a lower to higher apply status are multiplied by 0.61 (95% CI 0.39, 0.93) as we move from private to public, all else held constant.
- How about pared?

Comparison to a Null Model

```
mod_p0 <- polr(apply ~ 1, data = gradschool)

anova(mod_p1, mod_p0)
```

Likelihood ratio tests of ordinal regression models

Response: apply

	Model	Resid. df	Resid. Dev	Test	Df
1	1	528	975.1828		
2	pared + public + gpa	525	900.9629	1 vs 2	3
	LR stat.	Pr(Chi)			
1					
2	74.21989	5.551115e-16			

AIC and BIC are available, too

We could also compare model `mod_p1` to the null model `mod_p0` with AIC or BIC.

```
AIC(mod_p1, mod_p0)
```

	df	AIC
mod_p1	5	910.9629
mod_p0	2	979.1828

```
BIC(mod_p1, mod_p0)
```

	df	BIC
mod_p1	5	932.3273
mod_p0	2	987.7286

Testing the Proportional Odds Assumption

One way to test the proportional odds assumption is to compare the fit of the proportional odds logistic regression to a model that does not make that assumption. A natural candidate is a **multinomial logit** model, which is typically used to model unordered multi-categorical outcomes, and fits a slope to each level of the `apply` outcome in this case, as opposed to the proportional odds logit, which fits only one slope across all levels.

Since the proportional odds logistic regression model is nested in the multinomial logit, we can perform a likelihood ratio test. To do this, we first fit the multinomial logit model, with the `multinom` function from the `nnet` package.

Fitting the multinomial model

Again, the multinomial model is fit here using `multinom()` from the **nnet** package...

```
m1_multi <- multinom(apply ~ pared + public + gpa,  
                      data = gradschool)
```

```
# weights: 15 (8 variable)  
initial  value 582.264513  
iter 10 value 446.199617  
final   value 445.443366  
converged
```


The multinomial model

```
m1_multi
```

Call:

```
multinom(formula = apply ~ pared + public + gpa, data = gradsc
```

Coefficients:

	(Intercept)	pared	public	gpa
somewhat likely	-3.527249	1.072451	-0.97765580	0.9857488
very likely	-7.311227	1.400955	-0.02934361	1.6937996

Residual Deviance: 890.8867

AIC: 906.8867

Comparing the Models

The multinomial logit fits two intercepts and six slopes, for a total of 8 estimated parameters.

The proportional odds logit, as we've seen, fits two intercepts and three slopes, for a total of 5. The difference is 3, and we use that number in the sequence below to build our test of the proportional odds assumption.

Testing the Proportional Odds Assumption

```
LL_1 <- logLik(mod_p1)
LL_1m <- logLik(m1_multi)
(G <- -2 * (LL_1[1] - LL_1m[1]))
```

```
[1] 10.07618
```

```
pchisq(G, 3, lower.tail = FALSE)
```

```
[1] 0.01792959
```

The p value is 0.018, so it indicates that the proportional odds model fits less well than the more complex multinomial logit.

Comparing AIC and BIC

```
AIC(mod_p1)
```

```
[1] 910.9629
```

```
AIC(m1_multi)
```

```
[1] 906.8867
```

```
BIC(mod_p1)
```

```
[1] 932.3273
```

```
BIC(m1_multi)
```

```
[1] 941.0697
```

What to do in light of these results...

- A non-significant p value here isn't always the best way to assess the proportional odds assumption, but it does provide some evidence of model adequacy.
- The stronger BIC (and only slightly worse AIC) for our POLR model relative to the multinomial gives us some conflicting advice.
 - One alternative would be to fit the multinomial model instead.
 - Another would be to fit a check of residuals (see Frank Harrell's RMS text.)
 - Another would be to fit a different model for ordinal regression. Several are available (check out `orm` in the `rms` package, for instance.)

Section 3

Using `lrm` for Proportional Odds Logistic Regression

Using lrm to work through this model

```
d <- datadist(gradschool)
options(datadist = "d")
mod <- lrm(apply ~ pared + public + gpa,
          data = gradschool, x = T, y = T)
```

mod output

```
> mod
Logistic Regression Model

1rm(formula = apply ~ pared + public + gpa, data = gradschool,
     x = T, y = T)
```

		Model Likelihood	Discrimination	Rank Discrim.
		Ratio Test	Indexes	Indexes
obs	530	LR chi2 74.22	R2 0.155	C 0.684
unlikely	303	d.f. 3	g 0.895	Dxy 0.369
somewhat likely	172	Pr(> chi2) <0.0001	gr 2.448	gamma 0.369
very likely	55		gp 0.200	tau-a 0.206
max deriv	5e-09		Brier 0.216	

	Coef	S.E.	Wald Z	Pr(> Z)
y>=somewhat likely	-3.8728	0.5721	-6.77	<0.0001
y>=very likely	-5.9413	0.6063	-9.80	<0.0001
pared	1.1525	0.2184	5.28	<0.0001
public	-0.4949	0.2195	-2.25	0.0242
gpa	1.1416	0.1850	6.17	<0.0001


summary(mod)

Effects

Response : apply

Factor	Low	High	Diff.	Effect	S.E.	Lower 0.95
pared	0.00	1.00	1.00	1.15250	0.21843	0.72436
Odds Ratio	0.00	1.00	1.00	3.16600	NA	2.06340
public	0.00	1.00	1.00	-0.49486	0.21951	-0.92509
Odds Ratio	0.00	1.00	1.00	0.60966	NA	0.39650
gpa	2.61	3.44	0.83	0.94756	0.15354	0.64662
Odds Ratio	2.61	3.44	0.83	2.57940	NA	1.90910
Upper 0.95						
1.580600						
4.857900						
-0.064629						
0.937410						
1.248500						
3.485100						

```
plot(summary(mod))
```



Coefficients in our equation

```
mod$coef
```


y>=somewhat likely	y>=very likely	pared
-3.872786	-5.941317	1.152479
public	gpa	
-0.494859	1.141633	

Nomogram of mod (code)

```
fun.1 <- function(x) 1 - plogis(x)
fun.3 <- function(x)
  plogis(x - mod$coef[1] + mod$coef[2])

plot(nomogram(mod,
  fun=list('Prob Y = 1 (unlikely)' = fun.1,
    'Prob Y = 3 (very likely)' = fun.3)))
```

Nomogram of mod (result)



```
set.seed(432); validate(mod)
```

	index.orig	training	test	optimism
Dxy	0.3687	0.3663	0.3646	0.0017
R2	0.1553	0.1528	0.1511	0.0018
Intercept	0.0000	0.0000	0.0231	-0.0231
Slope	1.0000	1.0000	1.0170	-0.0170
Emax	0.0000	0.0000	0.0078	0.0078
D	0.1382	0.1359	0.1340	0.0019
U	-0.0038	-0.0038	-0.4637	0.4599
Q	0.1419	0.1397	0.5978	-0.4581
B	0.2155	0.2136	0.2171	-0.0035
g	0.8954	0.8833	0.8814	0.0019
gp	0.2004	0.1958	0.1975	-0.0016

	index.corrected	n
Dxy	0.3670	40
R2	0.1536	40
Intercept	0.0231	40
Slope	1.0170	40

More to come.

We'll share another example of ordinal logistic regression before we're done.